

Manual && Reproduction Report for

Watch and Act: Multi-orientation Open-set Scene Text Recognition via Dynamic Expert Routing

Hardware requirements

GPU:

Testing: NVIDIA GPU with Turing or later structure and more than 4 Gib of vram.

(yes pascals, maxwells or even fermis may work, but you gonna need to know what you are doing)

Training: 24GiB GPU with Turing or later structure

(In theory, 16GiB ones work too but I am not promising that.)

CPU: X86 cpus with AVX2 instruction set

(CPUs without AVX2 may or may not work with different versions of torch, I don't know)

RAM: 16GiB (Testing) 32Gib (Training)

Disk: 500GiB (well just dig out one old hdd from your bucket, or ebay and that's it)

Internet: Yes

Software Requirements

1. Fresh installed Archlinux with these packages:

`sudo vim plasma-meta nvidia openssh firefox dolphin konsole wget git less`

2. The user name is set to lasercat and it has to be a sudoer (or you may need to go through the code to replace paths if something went south).

3. ***Keep important data off this device!***

I don't want to wipe your data due to one or two failed cd commands followed by mv and/or rm

Trivia: Manul, aka the Pallas' cat, is the oldest cat species still alive on the earth.



Environment setup

1. install the following packages once you boot in.

```
sudo pacman -Syu pycharm-community-edition plasma-meta python-paramiko python-lmdb  
python-numba python-opencv python-pillow python-pip python-pyqtgraph python-pytorch-opt-  
cuda python-regex python-scikit-learn python-scipy python-torchvision-cuda python-tqdm python-  
xmldict make gcc cmake unzip python-setuptools kate firefox
```

2. make dirs

```
mkdir ~/cat ~/ssddata ~/hydra_saves ~/ssddata/anchors ~/Downloads
```

3. setup python stuff

```
mkdir ~/catvenv/; python3 -m venv ${PWD}/catvenv --system-site-packages;  
source ~/catvenv/bin/activate; pip install easydict editdistance wandb  
cd ~/cat; git clone https://github.com/lancercat/make_envNG.git  
cd make_envNG/ ; sh pylcs.sh ;  
unzip pytorch_scatter-laser.zip ; cd pytorch_scatter-2.1.2/; python setup.py install
```

4. download the data


```
cd ~/Download;  
wget https://www.kaggle.com/api/v1/datasets/download/object300/mose-extra;  
wget https://www.kaggle.com/api/v1/datasets/download/vsdf2898kaggle/osocrtraining;
```

5. download models to ~/Downloads: 310-rel.tgz from

<https://www.kaggle.com/models/object300/wna-zero>

6. check the md5sum:

```
cd ~/Download; md5sum *;
```



```
[laser@cat@TESTMEOW ococr-kaggle]$ md5sum *  
12ff9bf65774f9c85066feebd683f1a4 310-rel.tgz  
d4246f8296151e1b996aed5bbb4f7e9d mose-extra  
a63cae694dd5c5bb4dbe2995bc811276 osocrtraining  
[laser@cat@TESTMEOW ococr-kaggle]$
```

```
12ff9bf65774f9c85066feebd683f1a4 310-rel.tgz  
d4246f8296151e1b996aed5bbb4f7e9d mose-extra  
a63cae694dd5c5bb4dbe2995bc811276 osocrtraining
```

Note if you have different md5, it could be normal or not--- as we can have slightly corrupted archives as well --- they are pulled from kaggle too but at least ours work :)

we pulled them from kaggle to make sure the online version do work

In short, if you encounter problems AND have different md5sum --- try download again.

7. clone the code:

```
cd ~/cat; git clone https://github.com/lancercat/wna.git
```

8. run extract.sh:

```
cd ~/cat/wna/ ; sh extract.sh
```

⑨. stop and check:

```
[lancercat@TESTMEOW ~]$ ls ssddata/ hydra_saves/ cat
cat:
make_envNG wna

hydra_saves/:
aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_b32
aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_mjst_b32
aroute_nd_only-v6S-tiny-AAF
aroute_nd_only-v6S-tiny-AAF-01E
aroute_nd_only-v6S-tiny-AAF-01E-run2
aroute_nd_only-v6S-tiny-AAF-run2
aroute_nd_only-v6S-tiny-nedmix-AAF
aroute_nd_only-v6S-tiny-nedmix-AAF-ohem01
aroute_nd_only-v6S-tiny-nedmix-AAF-ohem01E
aroute_nd_only-v6S-tiny-nedmix-AAF-ohem01E-run2
aroute_nd_only-v6S-tiny-nedmix-AAF-run2
aroute_nd_only-v6S-tiny-nedmix-AAF-run2

pgroute_nd_only-v6S-tiny-AAF
pgroute_nd_only-v6S-tiny-AAF-01E
pgroute_nd_only-v6S-tiny-AAF-01E-run2
pgroute_nd_only-v6S-tiny-AAF-nedmix
pgroute_nd_only-v6S-tiny-AAF-nedmix-01E
pgroute_nd_only-v6S-tiny-AAF-nedmix-01E-run2
pgroute_nd_only-v6S-tiny-AAF-nedmix-run2
pgroute_nd_only-v6S-tiny-AAF-run2
rule_based-v6S-AAF
rule_based-v6S-AAF-ohem01E
rule_based-v6S-AAF-ohem01E-run2
rule_based-v6S-AAF-run2

ssddata/:
artdb_seen      ctwdb_seen      dicts      IC13_1015      lsvtdb_seen_NG  mltrchlat_seen_NG  mltrkr_hori      ST-abi
artdb_seen_NG   ctwdb_seen_NG   dictsv2     IIIT5k_3000    MJ-abi          mltrjp_hori        rctwtrdb_seen    SVT
athenaNG        CUTE80          IC03_867    lsvtdb_seen    mltrchlat_seen  mltrjp_hv         rctwtrdb_seen_NG

[lancercat@TESTMEOW ~]$
```



Reproducing ablative studies

1. run `ablative.py` (in a screen session as it takes some time):

```
python ablative.py 2>&1 | tee all.log
```

(our log is uploaded to github)

2. After the ablative experiments are finished (usually takes like a dozen hours on a GTX 1650), run `ablative2table.py` to compute mean and standard deviation:

```
/home/lasercat/catvenv/bin/python3.13 /home/lasercat/cat/wna/ablative2table.py
\newcommand{\LPAoBASEoJPNHVo6ZSLoLAoAVG}{40.97}
\newcommand{\LPAoBASEoJPNHVo6ZSLoLAoSTD}{1.02}
\newcommand{\LPAoW0ARoUeJPNHVo6ZSLoLAoAVG}{42.11}
\newcommand{\LPAoW0ARoUeJPNHVo6ZSLoLAoSTD}{0.51}
\newcommand{\LPAoPGoLoJPNHVo6ZSLoLAoAVG}{41.49}
\newcommand{\LPAoPGoLoJPNHVo6ZSLoLAoSTD}{1.01}
\newcommand{\LPAoPGoLoSoJPNHVo6ZSLoLAoAVG}{43.09}
\newcommand{\LPAoPGoLoSoJPNHVo6ZSLoLAoSTD}{0.33}
\newcommand{\LPAoPGoNoJPNHVo6ZSLoLAoAVG}{39.66}
\newcommand{\LPAoPGoNoJPNHVo6ZSLoLAoSTD}{0.36}
\newcommand{\LPAoPGoNoJPNHVo6ZSLoLAoSTD}{0.36}
\newcommand{\LPAoPGoNoSoJPNHVo6ZSLoLAoAVG}{43.17}
\newcommand{\LPAoPGoNoSoJPNHVo6ZSLoLAoSTD}{0.39}
\newcommand{\LPAoARoLoJPNHVo6ZSLoLAoAVG}{41.88}
\newcommand{\LPAoARoLoJPNHVo6ZSLoLAoSTD}{0.59}
\newcommand{\LPAoARoLoSoJPNHVo6ZSLoLAoAVG}{43.90}
\newcommand{\LPAoARoLoSoJPNHVo6ZSLoLAoSTD}{0.86}
\newcommand{\LPAoW0oHEMoJPNHVo6ZSLoLAoAVG}{42.05}
\newcommand{\LPAoW0oHEMoJPNHVo6ZSLoLAoSTD}{0.25}
\newcommand{\LPAoJPNHVo6ZSLoLAoAVG}{45.14}
\newcommand{\LPAoJPNHVo6ZSLoLAoSTD}{0.35}
\newcommand{\LPAoOHEMoJPNHVo6ZSLoLAoAVG}{44.53}
\newcommand{\LPAoOHEMoJPNHVo6ZSLoLAoSTD}{0.31}
-----
Process finished with exit code 0
```

3. save the output text as a txt file, say `reprod.txt`

(our log can be found in this github repo as `reprod.txt`)

4. diff them:

```
diff reprod.txt paperref.txt --color
```

Expect the results to be a bit different here and there due to float errors between different implementations of operators, which can be caused by different torch versions, gpu structures, vram, etc (yes, torch[1] and CUDNN[2] may choose different backend on different hardware)

[1] <https://discuss.pytorch.org/t/different-machines-different-results/100126>

[2] <https://docs.nvidia.com/deeplearning/cudnn/backend/latest/developer/core-concepts.html>

```
(catvenv) [lasercat@TESTMEOW wna]$ diff reprod.txt paperref.txt --color
2c2
< \newcommand{\LPAoBASEoJPNHVo6ZSLoLAoSTD}{1.02}
---
> \newcommand{\LPAoBASEoJPNHVo6ZSLoLAoSTD}{1.03}
7c7
< \newcommand{\LPAoPGoLoSoJPNHVo6ZSLoLAoAVG}{43.09}
---
> \newcommand{\LPAoPGoLoSoJPNHVo6ZSLoLAoAVG}{43.07}
9,10c9,10
< \newcommand{\LPAoPGoNoJPNHVo6ZSLoLAoAVG}{39.66}
< \newcommand{\LPAoPGoNoJPNHVo6ZSLoLAoSTD}{0.36}
---
> \newcommand{\LPAoPGoNoJPNHVo6ZSLoLAoAVG}{39.65}
> \newcommand{\LPAoPGoNoJPNHVo6ZSLoLAoSTD}{0.35}
13c13
< \newcommand{\LPAoARoLoJPNHVo6ZSLoLAoAVG}{41.88}
---
> \newcommand{\LPAoARoLoJPNHVo6ZSLoLAoAVG}{41.89}
18,20c18,20
< \newcommand{\LPAoW0oHEMoJPNHVo6ZSLoLAoSTD}{0.25}
< \newcommand{\LPAoJPNHVo6ZSLoLAoAVG}{45.14}
< \newcommand{\LPAoJPNHVo6ZSLoLAoSTD}{0.35}
---
> \newcommand{\LPAoW0oHEMoJPNHVo6ZSLoLAoSTD}{0.24}
> \newcommand{\LPAoJPNHVo6ZSLoLAoAVG}{45.15}
> \newcommand{\LPAoJPNHVo6ZSLoLAoSTD}{0.36}
22a23,24
```



Reproducing Regular Model


run object310-rel/project310_v6SF_stability-re/aroute_nd_only-v6S-tiny-nedmix-AAF-ohem01E

```
INFO: starting
/home/lasercat/cat/ana/conda/envs/conda-framework-NG/modules/concat_mvn_dev.py:67: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.detach().clone() or sourceTensor.detach().clone().requires_grad_(True)
  img=torch.stack([torch.tensor(i, dtype=this.get_type()) for i in imgelist]).permute(0,3,1,2).contiguous().to(this.mean.data.device)-this.mean;
WARN: Corrupted image for 5171
Date: 2025-06-12 14:36:05.161692 ,TEST: KR-KR-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 5170 ,ACR: 0.195357833655706 ,Lenpred_ACR: 0.6969052224371374 ,FPS: 107.1092715124331
WARN: Corrupted image for 5075
Date: 2025-06-12 14:36:52.248106 ,TEST: JPNHV-JPNHV-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 5074 ,ACR: 0.45309420575482856 ,Lenpred_ACR: 0.8281434765471029 ,FPS: 107.76280253973862
WARN: Corrupted image for 5075
{'Date': datetime.datetime(2025, 6, 12, 14, 37, 38, 995094), 'TEST': 'JPNHV-JPNHV-OSR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.7602158028064765, 'R': 0.7063807254434736, 'P': 0.9467908961593172, 'F': 0.8098345727728031}
WARN: Corrupted image for 5075
{'Date': datetime.datetime(2025, 6, 12, 14, 38, 26, 30389), 'TEST': 'JPNHV-JPNHV-GOSR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.687875574407918, 'R': 0.6534521158129176, 'P': 0.8634490876986463, 'F': 0.7439148073022313}
WARN: Corrupted image for 5075
{'Date': datetime.datetime(2025, 6, 12, 14, 39, 13, 820), 'TEST': 'JPNHV-JPNHV-OSTR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.7174177831912302, 'R': 0.802757564151666, 'P': 0.9109083007388092, 'F': 0.8534201954397395}
Date: 2025-06-12 14:39:46.950327 ,TEST: JPN-JPN-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 4009 ,ACR: 0.4784235470192068 ,Lenpred_ACR: 0.8233973559491145 ,FPS: 118.08836023802422
{'Date': datetime.datetime(2025, 6, 12, 14, 40, 20, 974879), 'TEST': 'JPN-JPN-OSR', 'Epoch': 0, 'Iter': 0, 'Total': 4009.0, 'KACR': 0.7983433981385729, 'R': 0.715318869165023, 'P': 0.958590308370044, 'F': 0.8192771084337348}
{'Date': datetime.datetime(2025, 6, 12, 14, 40, 55, 91256), 'TEST': 'JPN-JPN-GOSR', 'Epoch': 0, 'Iter': 0, 'Total': 4009.0, 'KACR': 0.7238437359676695, 'R': 0.6604938271604939, 'P': 0.8923426838514026, 'F': 0.7591899645275717}
{'Date': datetime.datetime(2025, 6, 12, 14, 41, 28, 923666), 'TEST': 'JPN-JPN-OSTR', 'Epoch': 0, 'Iter': 0, 'Total': 4009.0, 'KACR': 0.7531380753138075, 'R': 0.8235574630424416, 'P': 0.9421713038734315, 'F': 0.878880407124682}
wandb:
wandb: You can sync this run to the cloud by running:
wandb: wandb sync /home/lasercat/hydra_logs/aroute_nd_only-v6S-tiny-nedmix-AAF-ohem01E/watch_and_control/aroute_nd_only-v6S-tiny-nedmix-AAF-ohem01E/wandb/offline-run-20250612_143501-93thcnmc
wandb: Find logs at: ../../../../hydra_logs/aroute_nd_only-v6S-tiny-nedmix-AAF-ohem01E/watch_and_control/aroute_nd_only-v6S-tiny-nedmix-AAF-ohem01E/wandb/offline-run-20250612_143501-93thcnmc/logs
Process finished with exit code 0
```

Let's map them to the table

Split	Registered	Out-of-Set	Name	LA	Recall	Precision	FM
GZSL (Hori.)	<i>Unique Kanji</i>	∅	OSOCR-Large [4]	30.83	—	—	—
	Shared Kanji		OpenCCD-Large [23]	41.31	—	—	—
	<i>Kana</i> , Latin,		OpenSAVR-XL [34]	42.58	—	—	—
			MOoSE-XL* [7]	39.56	—	—	—
			CFOR-XL [37]	44.47	—	—	—
			WnA-S(Ours)* ---	47.89	—	—	—
Date: 2025-06-12 14:39:46.950327 ,TEST: JPN-JPN-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 4009 ,ACR: 0.4784235470192068 ,Lenpred_ACR: 0.8233973559491145 ,FPS: 118.08836023802422							
OSTR (Hori.)	Shared Kanji	<i>Kana</i> Latin	OSOCR-Large [4]	58.57	24.46	93.78	38.80
	<i>Unique Kanji</i>		OpenSAVR-XL [34]	72.33	72.96	92.62	81.62
			MOoSE-XL* [7]	64.80	80.49	89.12	84.50
			CFOR [37]	71.80	86.36	89.52	87.91
			WnA-S(Ours)* ---	75.31	82.30	94.21	87.85
			s), 'TEST': 'JPN-JPN-OSTR', 'Epoch': 0, 'Iter': 0, 'Total': 4009.0, 'KACR': 0.7531380753138075, 'R': 0.8235574630424416, 'P': 0.9421713038734315, 'F': 0.878880407124682}				
GZSL (MO)	Shared Kanji	∅	MOoSE-XL	37.39	—	—	—
	<i>Unique Kanji</i>		WnA-S(Ours) ---	45.35	—	—	—
Date: 2025-06-12 14:36:52.248106 ,TEST: JPNHV-JPNHV-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 5074 ,ACR: 0.45309420575482856 ,Lenpred_ACR: 0.8281434765471029 ,FPS: 107.76280253973862							
OSR (MO)	Shared Kanji	<i>Unique Kanji</i> Latin	MOoSE-XL	75.56	74.05	95.79	83.53
			WnA-S(Ours) ---	76.02	70.63	94.87	80.98
'TEST': 'JPNHV-JPNHV-OSR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.7602158828064765, 'R': 0.7063807254434736, 'P': 0.9467908961593172, 'F': 0.8098345727728031}							
GOSR (MO)	Shared Kanji	<i>Kana</i>	MOoSE-XL	59.81	63.83	81.89	71.74
	<i>Unique Kanji</i>		WnA-S(Ours) ---	68.82	65.39	86.35	74.42
'TEST': 'JPNHV-JPNHV-GOSR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.687875574407918, 'R': 0.6534521158129176, 'P': 0.8634490876986463, 'F': 0.7439148073022313}							
OSTR (MO)	Shared Kanji	Latin	MOoSE-XL	61.75	80.01	88.18	83.90
	<i>Unique Kanji</i>		WnA-S(Ours) ---	71.78	80.28	91.09	85.34
'TEST': 'JPNHV-JPNHV-OSTR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.7174177831912302, 'R': 0.802757564151666, 'P': 0.9109083007388092, 'F': 0.8534201954397395}							
GZSL (KR)	<i>Hangul</i>	∅	CFOR-XL [37]	22.14	—	—	—
	Latin		WnA-S(Ours) ---	19.52	—	—	—
Date: 2025-06-12 14:36:05.161692 ,TEST: KR-KR-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 5170 ,ACR: 0.195357833655706 ,Lenpred_ACR: 0.6969052224371374 ,FPS: 107.1092715124331							

Note the performance differs a bit on this 1650, but not much (utilization btw)



Reproducing Large Model

Open-set

Run object310-rel/XL/aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_b32/osocr_benchall.py

```
/home/lasercat/cat/ana/ana_env/ana_framework_00/modules/concat_env_dev.py:67: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.detach().clone() or sourceTensor.detach().clone().requires_grad_(True), rather than torch.tensor(sourceTensor).
  img=torch.stack([torch.tensor(i, dtype=this.get_type()) for i in imgList]).permute(0,3,1,2).contiguous().to(this.mean_data.device)-this.mean;
WARN: Corrupted image for 5871
Date: 2025-06-12 16:06:15.56954 ,TEST: KR-KR-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 5170 ,ACR: 0.23965183752417796 ,Lenpred_ACR: 0.7253384912959381 ,FPS: 93.31770637168638
WARN: Corrupted image for 5875
Date: 2025-06-12 16:07:21.041607 ,TEST: JPNHV-JPNHV-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 5074 ,ACR: 0.461568782026015 ,Lenpred_ACR: 0.8303113914071738 ,FPS: 91.47280815247775
WARN: Corrupted image for 5875
('Date': datetime.datetime(2025, 6, 12, 16, 0, 40, 4023), 'TEST': 'JPNHV-JPNHV-OSR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.7764070932922128, 'R': 0.7151178183743712, 'P': 0.954416961138742, 'F': 0.8176176782208696)
WARN: Corrupted image for 5875
('Date': datetime.datetime(2025, 6, 12, 16, 9, 11, 852653), 'TEST': 'JPNHV-JPNHV-GOSR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.7122658183103571, 'R': 0.6298440979955456, 'P': 0.8815461346633416, 'F': 0.7347362951415952)
WARN: Corrupted image for 5875
('Date': datetime.datetime(2025, 6, 12, 16, 9, 386266), 'TEST': 'JPNHV-JPNHV-OSTR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.738936256976452, 'R': 0.8050555342780544, 'P': 0.9239560439560439, 'F': 0.860417519433075)
Date: 2025-06-12 16:10:48.242145 ,TEST: JPN-JPN-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 4809 ,ACR: 0.4804190571214767 ,Lenpred_ACR: 0.823896233474682 ,FPS: 98.4153275062363
('Date': datetime.datetime(2025, 6, 12, 16, 11, 28, 836917), 'TEST': 'JPN-JPN-OSR', 'Epoch': 0, 'Iter': 0, 'Total': 4809.0, 'KACR': 0.809720785938841, 'R': 0.7176199868507561, 'P': 0.9620978404858317, 'F': 0.8228674072679366)
('Date': datetime.datetime(2025, 6, 12, 16, 12, 9, 855949), 'TEST': 'JPN-JPN-GOSR', 'Epoch': 0, 'Iter': 0, 'Total': 4809.0, 'KACR': 0.763152227213496, 'R': 0.6254567981234568, 'P': 0.9076797385620915, 'F': 0.7391882908866937)
('Date': datetime.datetime(2025, 6, 12, 16, 12, 50, 407182), 'TEST': 'JPN-JPN-OSTR', 'Epoch': 0, 'Iter': 0, 'Total': 4809.0, 'KACR': 0.7335355648535565, 'R': 0.806390810681927, 'P': 0.9489337822671156, 'F': 0.8718741943799947)
wandb:
wandb: You can sync this run to the cloud by running:
wandb: wandb sync /home/lasercat/hydra_logs/aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_b32/watch_and_control/aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_b32/wandb/offline-run-20250612_160516-78pek5ea
wandb: Find logs at: .../hydra_logs/aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_b32/watch_and_control/aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_b32/wandb/offline-run-20250612_160516-78pek5ea/logs
Process finished with exit code 0
```

Let's map them to the table

Split	Registered	Out-of-Set	Name	LA	Recall	Precision	FM
GZSL (Hori.)	Unique Kanji	∅	OSOCR-Large [4]	30.83	—	—	—
	Shared Kanji,		OpenCCD-Large [23]	41.31	—	—	—
	Kana, Latin,		OpenSAVR-XL [34]	42.58	—	—	—
			MOoSE-XL* [7]	39.56	—	—	—
			CFOR-XL [37]	44.47	—	—	—

Date: 2025-06-12 16:10:48.242145 ,TEST: JPN-JPN-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 4809 ,ACR: 0.4804190571214767 ,Lenpred_ACR: 0.823896233474682 ,FPS: 98.4153275062363

			WnA-S-XL (Ours)*	48.02	—	—	—
OSTR (Hori.)	Shared Kanji,	Kana Latin	OSOCR-Large [4]	58.57	24.46	93.78	38.80
	Unique Kanji		OpenSAVR-XL [34]	72.33	72.96	92.62	81.62
			MOoSE-XL* [7]	64.80	80.49	89.12	84.50
			CFOR [37]	71.80	86.36	89.52	87.91

'TEST': 'JPN-JPN-OSTR', 'Epoch': 0, 'Iter': 0, 'Total': 4809.0, 'KACR': 0.7735355648535565, 'R': 0.806390810681927, 'P': 0.9489337822671156, 'F': 0.8718741943799947

			WnA-S-XL(Ours)*	77.35	80.68	94.89	87.21
--	--	--	-----------------	-------	-------	-------	-------

GZSL	Shared Kanji		MOoSE-XL	37.39	—	—	—
------	--------------	--	----------	-------	---	---	---

: 2025-06-12 16:07:21.041607 ,TEST: JPNHV-JPNHV-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 5074 ,ACR: 0.461568782026015 ,Lenpred_ACR: 0.8303113914071738 ,FPS: 91.47280815247775

	Latin, Kana		WnA-S-XL(Ours)	46.14	—	—	—
--	-------------	--	----------------	-------	---	---	---

OSR	Shared Kanji	Unique Kanji	MOoSE-XL	75.56	74.05	95.79	83.53
-----	--------------	--------------	----------	-------	-------	-------	-------

ST': 'JPNHV-JPNHV-OSR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.7764070932922128, 'R': 0.7151178183743712, 'P': 0.954416961138742, 'F': 0.8176176782208696)

			WnA-S-XL(Ours)	77.64	71.45	95.44	81.72
--	--	--	----------------	-------	-------	-------	-------

GOSR	Shared Kanji	Kana	MOoSE-XL	59.81	63.83	81.89	71.74
------	--------------	------	----------	-------	-------	-------	-------

': 'JPNHV-JPNHV-GOSR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.7122658183103571, 'R': 0.6298440979955456, 'P': 0.8815461346633416, 'F': 0.7347362951415952)

	Latin		WnA-S-XL(Ours)	71.22	62.89	88.08	73.38
--	-------	--	----------------	-------	-------	-------	-------

OSTR	Shared Kanji	Latin	MOoSE-XL	61.75	80.01	88.18	83.90
------	--------------	-------	----------	-------	-------	-------	-------

'JPNHV-JPNHV-OSTR', 'Epoch': 0, 'Iter': 0, 'Total': 5074.0, 'KACR': 0.738936256976452, 'R': 0.8050555342780544, 'P': 0.9239560439560439, 'F': 0.860417519433075)


			WnA-S-XL(Ours)	73.89	80.50	92.39	86.04
--	--	--	----------------	-------	-------	-------	-------

GZSL	Hangul	∅	CFOR-XL [37]	22.14	—	—	—
------	--------	---	--------------	-------	---	---	---

TEST: KR-KR-GZSL ,Epoch: 0 ,Iter: 0 ,Total: 5170 ,ACR: 0.23965183752417796 ,Lenpred_ACR: 0.7253384912959381 ,FPS: 93.33770637168638

			WnA-S-XL (Ours)	24.00	—	—	—
--	--	--	-----------------	-------	---	---	---

Note the performance differs a bit on this 1650, but not much (utilization btw)



Reproducing Fig 1& Inference on your own data


Run object310-rel/XL/aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_b32/athena.py



And harvest the results under

object310-rel/XL/aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_b32/athenaNG_results

For inferencing your own data, you need to prepare the meta/alphabets.txt, the meta/notofont.ttf, and images.



The meta/alphabets.txt contains all alphabets, each a line. If an alphabet has more than one case, split them with space.

For the notofont.ttf, you search for the font here: <https://fonts.google.com/noto>
Download it, extract the regular variant, rename it to notofont.ttf, and add it to the folder.

Close-set

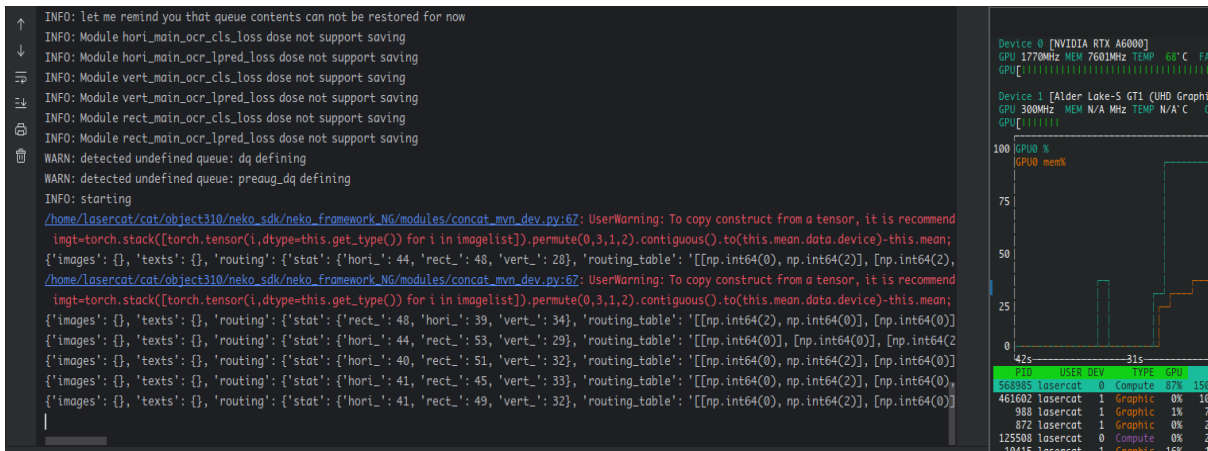
Run object310-rel/XL/aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_mjst_b32/test.py




Training from Scratch

Run train.py under the method folder you choose. For example to train the large model:

object310-rel/XL/aroute_nd_only-v6S-reg-nedmix-AAF-ohem01E_promelas_b32/train.py



And watch it cook. Utilization BTW



If you want to train with your own data, please try to follow the instructions from the vanilla osocr-data repo, the lmdbs are fully compatible.

The newly added dicts_v2/lang/vismeta.pt builders can be reverse engineered from the athena scripts. Writing this manual takes too long already so I am not going to write an extensive document on this now... GLHF.



That's it

If anything is buggy, doesn't work, or you just want to chat, open an issue or contact us

Chang: lasercat@gmx.us (all time); chang.liu@ltu.se (work time please)

Elisa: elisa.barney@ltu.se

BTW... Since “manul” is spelled like “manual”, so you will see a manul at the end of each section as an end-of-section icon.

See ya && GLHF

